

Ensemble Methods

Zander Bonnet

September 25, 2024

```
In [59]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
plt.style.use('default')
import seaborn as sns
from sklearn.ensemble import GradientBoostingRegressor, BaggingRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
import xgboost as xgb
from catboost import CatBoostRegressor
from sklearn.tree import DecisionTreeRegressor
```

```
In [60]: from ucimlrepo import fetch_ucirepo

# fetch dataset
bike_sharing = fetch_ucirepo(id=275)

# data (as pandas dataframes)
X = bike_sharing.data.features
y = bike_sharing.data.targets

# metadata
print(bike_sharing.metadata)

# variable information
print(bike_sharing.variables)
```

```
{'uci_id': 275, 'name': 'Bike Sharing', 'repository_url': 'https://archive.i
cs.uci.edu/dataset/275/bike+sharing+dataset', 'data_url': 'https://archive.i
cs.uci.edu/static/public/275/data.csv', 'abstract': 'This dataset contains t
he hourly and daily count of rental bikes between years 2011 and 2012 in Cap
ital bikeshare system with the corresponding weather and seasonal informatio
n.', 'area': 'Social Science', 'tasks': ['Regression'], 'characteristics':
['Multivariate'], 'num_instances': 17389, 'num_features': 13, 'feature_type
s': ['Integer', 'Real'], 'demographics': [], 'target_col': ['cnt'], 'index_c
ol': ['instant'], 'has_missing_values': 'no', 'missing_values_symbol': None,
'year_of_dataset_creation': 2013, 'last_updated': 'Sun Mar 10 2024', 'databse
_t_doi': '10.24432/C5W894', 'creators': ['Hadi Fanaee-T'], 'intro_paper': {'I
D': 422, 'type': 'NATIVE', 'title': 'Event labeling combining ensemble detec
tors and background knowledge', 'authors': 'Hadi Fanaee-T, João Gama', 'venu
e': 'Progress in Artificial Intelligence', 'year': 2013, 'journal': None, 'D
OI': '10.1007/s13748-013-0040-3', 'URL': 'https://www.semanticscholar.org/pa
per/bc42899f599d31a5d759f3e0a3ea8b52479d6423', 'sha': None, 'corpus': None,
'arxiv': None, 'mag': None, 'acl': None, 'pmid': None, 'pmcid': None}, 'addi
tional_info': {'summary': 'Bike sharing systems are new generation of tradit
ional bike rentals where whole process from membership, rental and return ba
ck has become automatic. Through these systems, user is able to easily rent
a bike from a particular position and return back at another position. Curre
ntly, there are about over 500 bike-sharing programs around the world which
is composed of over 500 thousands bicycles. Today, there exists great intere
st in these systems due to their important role in traffic, environmental an
d health issues. \r\n\r\nApart from interesting real world applications of b
ike sharing systems, the characteristics of data being generated by these sy
stems make them attractive for the research. Opposed to other transport serv
ices such as bus or subway, the duration of travel, departure and arrival po
sition is explicitly recorded in these systems. This feature turns bike shar
ing system into a virtual sensor network that can be used for sensing mobili
ty in the city. Hence, it is expected that most of important events in the c
ity could be detected via monitoring these data.', 'purpose': None, 'funded_
by': None, 'instances_represent': None, 'recommended_data_splits': None, 'se
nsitive_data': None, 'preprocessing_description': None, 'variable_info': 'Bo
th hour.csv and day.csv have the following fields, except hr which is not av
ailable in day.csv\r\n\r\n\t\r\n\t\t- instant: record index\r\n\r\n\t\t- dteday : date\r
\r\n\t\t- season : season (1:winter, 2:spring, 3:summer, 4:fall)\r\n\r\n\t\t- yr : yea
r (0: 2011, 1:2012)\r\n\r\n\t\t- mnth : month ( 1 to 12)\r\n\r\n\t\t- hr : hour (0 to 2
3)\r\n\r\n\t\t- holiday : weather day is holiday or not (extracted from http://dch
r.dc.gov/page/holiday-schedule)\r\n\r\n\t\t- weekday : day of the week\r\n\r\n\t\t- work
ingday : if day is neither weekend nor holiday is 1, otherwise is 0.\r\n\r\n\t\t+
weathersit : \r\n\r\n\t\t\t- 1: Clear, Few clouds, Partly cloudy, Partly cloudy\r
\r\n\t\t\t- 2: Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist\r\n\r
n\t\t\t- 3: Light Snow, Light Rain + Thunderstorm + Scattered clouds, Light Rai
n + Scattered clouds\r\n\r\n\t\t\t- 4: Heavy Rain + Ice Pallets + Thunderstorm + M
ist, Snow + Fog\r\n\r\n\t\t- temp : Normalized temperature in Celsius. The values
are derived via (t-t_min)/(t_max-t_min), t_min=-8, t_max=+39 (only in hourly
scale)\r\n\r\n\t\t- atemp: Normalized feeling temperature in Celsius. The values a
re derived via (t-t_min)/(t_max-t_min), t_min=-16, t_max=+50 (only in hourly
scale)\r\n\r\n\t\t- hum: Normalized humidity. The values are divided to 100 (max)
\r\n\r\n\t\t- windspeed: Normalized wind speed. The values are divided to 67 (max)
\r\n\r\n\t\t- casual: count of casual users\r\n\r\n\t\t- registered: count of registered
users\r\n\r\n\t\t- cnt: count of total rental bikes including both casual and regi
stered\r\n\r\n', 'citation': None}}
```

	name	role	type	demographic	\
0	instant	ID	Integer	None	

1	dteday	Feature	Date	None
2	season	Feature	Categorical	None
3	yr	Feature	Categorical	None
4	mnth	Feature	Categorical	None
5	hr	Feature	Categorical	None
6	holiday	Feature	Binary	None
7	weekday	Feature	Categorical	None
8	workingday	Feature	Binary	None
9	weathersit	Feature	Categorical	None
10	temp	Feature	Continuous	None
11	atemp	Feature	Continuous	None
12	hum	Feature	Continuous	None
13	windspeed	Feature	Continuous	None
14	casual	Other	Integer	None
15	registered	Other	Integer	None
16	cnt	Target	Integer	None

		description	units	missing_values
0		record index	None	no
1		date	None	no
2		1:winter, 2:spring, 3:summer, 4:fall	None	no
3		year (0: 2011, 1: 2012)	None	no
4		month (1 to 12)	None	no
5		hour (0 to 23)	None	no
6	weather	day is holiday or not (extracted from ...	None	no
7		day of the week	None	no
8	if	day is neither weekend nor holiday is 1, ot...	None	no
9	- 1:	Clear, Few clouds, Partly cloudy, Partly ...	None	no
10	Normalized	temperature in Celsius. The values ...	C	no
11	Normalized	feeling temperature in Celsius. The...	C	no
12	Normalized	humidity. The values are divided to...	None	no
13	Normalized	wind speed. The values are divided ...	None	no
14		count of casual users	None	no
15		count of registered users	None	no
16	count	of total rental bikes including both cas...	None	no

This data set is a collection of data about the amount of rental bikes being used. The dataset contains data such as the time of day, the type of weather, the day of the week, and so on. The description of the complete data set is printed above. With this data, I hope to accurately predict the number of rental bikes being used at a given time.

I will utilize various ensemble methods to find the most effective approach to accomplish this.

How accurately can I predict the number of rentals at a given time based on the factors in the dataset?

What ensemble method will prove to be the most accurate and efficient?

```
In [61]: #drop unwanted column
X = X.drop('dteday', axis = 1)
X.dtypes
```

```
Out[61]: season      int64
yr      int64
mnth     int64
hr      int64
holiday  int64
weekday  int64
workingday int64
weathersit int64
temp     float64
atemp    float64
hum      float64
windspeed float64
dtype: object
```

```
In [62]: #cast the categorical variables to categorical data type
cat = ['season', 'yr', 'mnth', 'hr', 'holiday', 'weekday', 'workingday', 'we
for col in cat:
    X[col] = X[col].astype('category', copy=False)
X.dtypes
```

```
Out[62]: season      category
yr      category
mnth     category
hr      category
holiday  category
weekday  category
workingday category
weathersit category
temp     float64
atemp    float64
hum      float64
windspeed float64
dtype: object
```

```
In [63]: X.head()
```

```
Out[63]:
```

	season	yr	mnth	hr	holiday	weekday	workingday	weathersit	temp	atemp	h
0	1	0	1	0	0	6	0	1	0.24	0.2879	C
1	1	0	1	1	0	6	0	1	0.22	0.2727	0
2	1	0	1	2	0	6	0	1	0.22	0.2727	0
3	1	0	1	3	0	6	0	1	0.24	0.2879	0
4	1	0	1	4	0	6	0	1	0.24	0.2879	0

```
In [64]: X.describe()
```

Out [64]:

	temp	atemp	hum	windspeed
count	17379.000000	17379.000000	17379.000000	17379.000000
mean	0.496987	0.475775	0.627229	0.190098
std	0.192556	0.171850	0.192930	0.122340
min	0.020000	0.000000	0.000000	0.000000
25%	0.340000	0.333300	0.480000	0.104500
50%	0.500000	0.484800	0.630000	0.194000
75%	0.660000	0.621200	0.780000	0.253700
max	1.000000	1.000000	1.000000	0.850700

In [65]: `y.head()`

Out [65]:

	cnt
0	16
1	40
2	32
3	13
4	1

In [66]: `y.describe()`

Out [66]:

	cnt
count	17379.000000
mean	189.463088
std	181.387599
min	1.000000
25%	40.000000
50%	142.000000
75%	281.000000
max	977.000000

In [67]: `#split the data to a train and test set`
`X_train, X_test, y_train, y_test = train_test_split(`
`X, y, test_size=0.30)`

Base Decision Tree

This will serve as a baseline as many of these ensemble models utilize decision trees under the hood.


```
In [68]: mod = DecisionTreeRegressor()  
mod.fit(X_train, np.ravel(y_train))  
pred = mod.predict(X_test)  
print(np.sqrt(mean_squared_error(y_test, pred)))
```

57.900987948798196

```
In [69]: mod.score(X_test,y_test)
```

Out[69]: 0.897945498484886

```
In [70]: plt.scatter(y_test.index,np.ravel(y_test))  
plt.scatter(X_test.index, pred)  
plt.show()
```



This baseline model performs fairly well with an RMSE of 57.86 and an r^2 score of .897. This means the model is about 89% accurate and on average is 57.8 rentals off.

Bagging

Bagging works by bootstrapping the data into smaller subsets and running models on these individual subsets. These models all run at the same time. After the completion of the smaller models, the bagging algorithm averages the responses to create the final output.


```
In [71]: model0 = BaggingRegressor(DecisionTreeRegressor())
model0.fit(X_train, np.ravel(y_train))
pred_final0 = model0.predict(X_test)
print(np.sqrt(mean_squared_error(y_test, pred_final0)))
```

44.988698022038534

```
In [72]: model0.score(X_test,y_test)
```

Out[72]: 0.9383877502359166

```
In [73]: plt.scatter(y_test.index,np.ravel(y_test))
plt.scatter(X_test.index, pred_final0)
plt.show()
```



Comparing the bagging algorithm to the baseline we can see that there was an overall improvement in the ability to predict the data. The r^2_{score} increased to .936 and the RMSE is now as low as 45. The bagging algorithm was able to reduce the variation in the data by averaging over multiple iterations of models. This will mean that models that are heavily impacted by more extreme data points will be outweighed by the more accurate models.

Gradient Boosting

GBM is a boosting algorithm, meaning it sequentially creates models that build upon each other. To accomplish this the model creates a base model and then calculates the residuals to the target. There is then a subsequent model created to minimize the

residuals. Once a prediction is made it is added to the original prediction and the residuals are recalculated. This continues for the given number of iterations.


```
In [74]: model = GradientBoostingRegressor(n_estimators=500)
model.fit(X_train, np.ravel(y_train))
pred_final = model.predict(X_test)
print(np.sqrt(mean_squared_error(y_test, pred_final)))
```

51.79883825916279

```
In [75]: model.score(X_test, y_test)
```

Out[75]: 0.9183229099150216

```
In [76]: plt.scatter(y_test.index, np.ravel(y_test))
plt.scatter(X_test.index, pred_final)
plt.show()
```



```
In [77]: test_score = np.zeros(500, dtype=np.float64)
for i, y_pred in enumerate(model.staged_predict(X_test)):
    test_score[i] = np.sqrt(mean_squared_error(y_test, y_pred))


fig = plt.figure(figsize=(6, 6))
plt.subplot(1, 1, 1)
plt.title("Deviance")
plt.plot(
    np.arange(500) + 1,
    np.sqrt(model.train_score_),
    "b-",
    label="Training Set Deviance",
```



```

)
plt.plot(np.arange(500) + 1, test_score, "r-", label="Test Set Deviance")
plt.legend(loc="upper right")
plt.xlabel("Boosting Iterations")
plt.ylabel("RMSE")
fig.tight_layout()

```



This model performs better than the original model, but not as well as the bagging model. With an r_2 score of .918 and an RMSE of 51.8, the model is just slightly worse. Looking at the Deviance we can see that as the iterations pass the RMSE begins to flatten and stabilize. This is a sign that adding more iterations will begin to become less and less beneficial. This model can minimize the variation in predictions by focusing on reducing the residuals.

XGBoost

XGBoost is a direct improvement of the previous method. It can do this by introducing regularization to the data, tree pruning, and built-in cross-validation. These additions make it so the algorithm doesn't spend as much time on tree nodes that do not improve the performance of the model.


```
In [78]: model2=xgb.XGBRegressor(enable_categorical = True)
model2.fit(X_train, y_train, eval_set = [(X_test, y_test)], verbose=10)
pred_final2 = model2.predict(X_test)
print(np.sqrt(mean_squared_error(y_test, pred_final2)))
```

```
[0]    validation_0-rmse:138.60938
[10]    validation_0-rmse:48.91669
[20]    validation_0-rmse:43.90393
[30]    validation_0-rmse:42.72042
[40]    validation_0-rmse:42.12085
[50]    validation_0-rmse:42.02758
[60]    validation_0-rmse:41.74164
[70]    validation_0-rmse:41.62703
[80]    validation_0-rmse:41.59106
[90]    validation_0-rmse:41.44882
[99]    validation_0-rmse:41.44781
41.44781393322742
```


```
In [79]: model2.score(X_test,y_test)
```

```
Out[79]: 0.9477046132087708
```

```
In [80]: plt.scatter(y_test.index,np.ravel(y_test))
plt.scatter(X_test.index, pred_final2)
plt.show()
```



```
In [81]: plt.plot(range(100),model2.evals_result_['validation_0']['rmse'])  
plt.ylabel('RMSE')  
plt.xlabel('Iteration')  
plt.show()
```



This model performs the best so far with an RMSE of 41.3 and a r^2 score of .947. This shows that the model performs very well. Looking at the RMSE Vs. Iterations plot we can also see that the model converged very early. There was very minimal improvement in RMSE after about 25 iterations.

CatBoost

CatBoost is another boosting algorithm, but this one excels in datasets with large amounts of categorical features. It can do this by converting categorical factors to numerical values using a variety of statistical methods.

```
In [82]: model3=CatBoostRegressor(verbose = 50)
model3.fit(X_train, y_train, cat_features=np.where(X.dtypes == 'category')[0])
pred_final3 = model3.predict(X_test)
print(np.sqrt(mean_squared_error(y_test, pred_final3)))
```

Learning rate set to 0.07538

0:	learn: 172.9089452	test: 172.4738129	best: 172.4738129
(0)	total: 5.11ms	remaining: 5.1s	
50:	learn: 70.2833412	test: 61.6242359	best: 61.6242359 (5
0)	total: 186ms	remaining: 3.46s	
100:	learn: 61.2717337	test: 53.1365734	best: 53.1365734 (10
0)	total: 337ms	remaining: 3s	
150:	learn: 57.3629442	test: 50.3866016	best: 50.3866016 (15
0)	total: 521ms	remaining: 2.93s	
200:	learn: 54.8220014	test: 48.9598002	best: 48.9598002 (20
0)	total: 669ms	remaining: 2.66s	
250:	learn: 52.6467134	test: 47.7721806	best: 47.7721806 (25
0)	total: 834ms	remaining: 2.49s	
300:	learn: 50.9642809	test: 46.9484068	best: 46.9484068 (30
0)	total: 988ms	remaining: 2.29s	
350:	learn: 49.0840582	test: 46.0318651	best: 46.0318651 (35
0)	total: 1.13s	remaining: 2.08s	
400:	learn: 47.8470091	test: 45.4162815	best: 45.4162815 (40
0)	total: 1.26s	remaining: 1.88s	
450:	learn: 47.0441712	test: 45.0592163	best: 45.0592163 (45
0)	total: 1.38s	remaining: 1.68s	
500:	learn: 46.0683988	test: 44.7082869	best: 44.7082869 (50
0)	total: 1.51s	remaining: 1.51s	
550:	learn: 45.1930652	test: 44.2556508	best: 44.2556508 (55
0)	total: 1.66s	remaining: 1.35s	
600:	learn: 44.5278151	test: 43.9904833	best: 43.9904833 (60
0)	total: 1.79s	remaining: 1.19s	
650:	learn: 43.9830670	test: 43.7702176	best: 43.7702176 (65
0)	total: 1.91s	remaining: 1.02s	
700:	learn: 43.3379794	test: 43.5027646	best: 43.5027646 (70
0)	total: 2.03s	remaining: 866ms	
750:	learn: 42.7272985	test: 43.3653891	best: 43.3653891 (75
0)	total: 2.17s	remaining: 718ms	
800:	learn: 42.0503314	test: 43.1342819	best: 43.1342819 (80
0)	total: 2.29s	remaining: 569ms	
850:	learn: 41.2505193	test: 42.8686796	best: 42.8686796 (85
0)	total: 2.42s	remaining: 424ms	
900:	learn: 40.6203609	test: 42.6511590	best: 42.6511590 (90
0)	total: 2.55s	remaining: 280ms	
950:	learn: 40.1073898	test: 42.4619685	best: 42.4605777 (94
9)	total: 2.67s	remaining: 138ms	
999:	learn: 39.6903829	test: 42.3803977	best: 42.3803977 (99
9)	total: 2.79s	remaining: 0us	

bestTest = 42.38039774

bestIteration = 999


42.38039774008865

```
In [83]: model3.score(X_test,y_test)
```


```
Out[83]: 0.9453248133685026
```

```
In [84]: plt.scatter(y_test.index,np.ravel(y_test))  
plt.scatter(X_test.index, pred_final3)
```

```
plt.show()
```



```
In [85]: plt.plot(range(0,1000),model3.evals_result_['learn']['RMSE'], label = 'Train')
plt.plot(range(0,1000),model3.evals_result_['validation']['RMSE'], label = 'Validation')
plt.ylabel('RMSE')
plt.xlabel('Iterations')
plt.legend()
plt.show()
```



This method performs very similarly to XGBoost. It does perform just slightly worse with an RMSE of 42.4 and a r^2 score of .945. This model does take much longer to execute as well. Looking at the RMSE Vs. Iterations plot we can also see that the validation set converged much quicker than the training data. This shows that the model was good at predicting external data at a small number of iterations.

Overall, the best model was XGBoost as it provided the most accurate predictions, and ran relatively fast compared to similar performers.

We were able to answer both of our questions. We found that XGBoost is the most accurate and efficient method to model this dataset, and it has an accuracy of about 94.7%.

Reference

Fanaee-T, H. (2013). Bike Sharing [Dataset]. UCI Machine Learning Repository.
<https://doi.org/10.24432/C5W894>.